

Winter 18 – AMS206B Homework 3 Solution

1. Consider three independent random variables X_1, X_2 and X_3 such that $X_i \stackrel{iid}{\sim} \text{Gamma}(a_i, b)$. Let

$$\mathbf{Y} = (Y_1, Y_2, Y_3) = \left(\frac{X_1}{X_1 + X_2 + X_3}, \frac{X_2}{X_1 + X_2 + X_3}, \frac{X_3}{X_1 + X_2 + X_3} \right).$$

- (a) Show that $\mathbf{Y} \sim \text{Dirichlet}(a_1, a_2, a_3)$, a Dirichlet distribution.
 (b) How can this result be used to generate random variables according to a Dirichlet distribution? Write a simple function in **R** or **Matlab** (your choice) that takes as inputs n , the number of trivariate vectors to be generated, and $\mathbf{a} = (a_1, a_2, a_3)$ and generates as an output a matrix of size $n \times 3$ whose rows correspond to independent samples from a Dirichlet distribution with parameter (a_1, a_2, a_3) . (*Note:* the value of b is not important as long as the three X have the same value for b)

Solution:

- (a) First, we consider the joint density of the three independent Gamma-distributed RVs:

$$p(x_1, x_2, x_3) = \prod_{i=1}^3 p(x_i) = \prod_{i=1}^3 \frac{x_i^{a_i-1} e^{-x_i/b}}{\Gamma(a_i) b^{a_i}} = \frac{e^{\sum_{i=1}^3 x_i/b} \prod_{i=1}^3 x_i^{a_i-1}}{b^{\sum_{i=1}^3 a_i} \prod_{i=1}^3 \Gamma(a_i)}.$$

Since $Y_i = X_i / \sum_{i=1}^3 X_i$, we find

$$\begin{aligned} X_1 &= Y_1 Z \\ X_2 &= Y_2 Z \\ X_3 &= Y_3 Z = (1 - Y_1 - Y_2) Z, \end{aligned}$$

where $Z = \sum_{i=1}^3 X_i$. To obtain the joint distribution of (Y_1, Y_2, Z) , we find the Jacobian for this change of variables. The matrix is

$$J = \begin{pmatrix} \frac{dx_1}{dy_1} & \frac{dx_1}{dy_2} & \frac{dx_1}{dz} \\ \frac{dx_2}{dy_1} & \frac{dx_2}{dy_2} & \frac{dx_2}{dz} \\ \frac{dx_3}{dy_1} & \frac{dx_3}{dy_2} & \frac{dx_3}{dz} \end{pmatrix} = \begin{pmatrix} Z & 0 & Y_1 \\ 0 & Z & Y_2 \\ -Z & -Z & (1 - Y_1 - Y_2) \end{pmatrix}$$

So, the Jacobian, $|J|$ is Z^2 .

$$p(Y_1, Y_2, Z) = \frac{(y_1 z)^{a_1-1} e^{-y_1 z/b} (y_2 z)^{a_2-1} e^{-y_2 z/b} \{(1 - y_1 - y_2) z\}^{a_3-1} e^{-(1-y_1-y_2)z/b}}{b^{\sum_{i=1}^3 a_i} \prod_{i=1}^3 \Gamma(a_i)} z^2,$$

where $0 < y_1, y_2 < 1$, $y_1 + y_2 < 1$ and $0 < z$.

By letting $y_3 = 1 - y_1 - y_2$,

$$p(Y_1, Y_2, Z) = \frac{y_1^{a_1-1} y_2^{a_2-1} y_3^{a_3-1} z^{\sum_{i=1}^3 a_i-1} e^{z \sum_{i=1}^3 y_i/b}}{b^{\sum_{i=1}^3 a_i} \prod_{i=1}^3 \Gamma(a_i)}.$$

We now integrate out z to obtain $p(y_1, y_2, y_3)$.

$$\begin{aligned}
 p(y_1, y_2, y_3) &= \int_{\mathbb{R}} \frac{y_1^{a_1-1} y_2^{a_2-1} y_3^{a_3-1} \overbrace{z^{\sum_{i=1}^3 a_i - 1} e^{-z/b}}^{\text{kernel for Gamma}(\sum_{i=1}^3 a_i, b)}}{b^{\sum_{i=1}^3 a_i} \prod_{i=1}^3 \Gamma(a_i)} dz \\
 &= \frac{\Gamma(\sum_{i=1}^3 a_i)}{\prod_{i=1}^3 \Gamma(a_i)} y_1^{a_1-1} y_2^{a_2-1} y_3^{a_3-1},
 \end{aligned}$$

- (b) #a is a vector of length p; a=(a_1, a_2, ..., a_p)
 #n is the sample size

```

> dirichlet <- function(a, n){
  p <- length(a)
  y <- array(NA, dim=c(n, p)) #Each row of y is iid sample from Dir(a)

  for (i in 1:n) {
    tmp <- rgamma(p, a, 1)
    y[i, ] <- tmp / sum(tmp)
  }

  return(y)
}

```

2. Y follows an inverse Gamma distribution with shape parameter a and scale parameter b ($Y \sim \text{IG}(a, 1/b)$) if $Y = 1/X$ with $X \sim \text{Gamma}(a, b)$ (assume the Gamma distribution is parameterized so that $E(X) = ab$).

- (a) Find the density of Y .
 (b) Compute $E(Y^k)$. Do you need to impose any constrain on the problem for this expectation to exists?
 (c) Compare $E(Y^k)$ to $1/E(X^k)$ (hint: look at the ratio of the two quantities).

Solution:

- (a) Since $X \sim \text{Gamma}(a, b)$,

$$f(x) = \frac{1}{\Gamma(a)b^a} x^{a-1} \exp\left\{-\frac{x}{b}\right\}, \text{ for } x > 0.$$

Let $y = 1/x$. Then $x = 1/y$ and $\frac{dx}{dy} = -1/y^2$. Therefore,

$$p(y) = \frac{1}{\Gamma(a)b^a} \left(\frac{1}{y}\right)^{a-1} \exp\left\{-\frac{1}{by}\right\} \left|-\frac{1}{y^2}\right| = \frac{(1/b)^a}{\Gamma(a)} y^{-(a+1)} \exp\left\{-\frac{(1/b)}{y}\right\}$$

which is an inverse Gamma with shape a and rate $1/b$.

(b)

$$\begin{aligned} E(Y^k) &= \int_{\mathbb{R}^+} y^k \frac{(1/b)^a}{\Gamma(a)} y^{-(a+1)} \exp\left\{-\frac{(1/b)}{y}\right\} dy \\ &= \frac{(1/b)^a}{\Gamma(a)} \int_{\mathbb{R}^+} y^{-(a-k+1)} \exp\left\{-\frac{(1/b)}{y}\right\} dy = \frac{(1/b)^a}{\Gamma(a)} \frac{\Gamma(a-k)}{(1/b)^{a-k}} \end{aligned}$$

That is, $E(Y^k) = \Gamma(a-k)/(b^k \Gamma(a))$ for $a-k > 0$.

(c)

$$\begin{aligned} E(X^k) &= \int_{\mathbb{R}^+} x^k \frac{1}{\Gamma(a)b^a} x^{a-1} \exp\left\{-\frac{x}{b}\right\} dx \\ &= \frac{1}{\Gamma(a)b^a} \int_{\mathbb{R}^+} x^{a+k-1} \exp\left\{-\frac{x}{b}\right\} dx = \frac{1}{\Gamma(a)b^a} \Gamma(a+k) b^{a+k}. \end{aligned}$$

Then,

$$\zeta \equiv \frac{1/E(X^k)}{E(Y^k)} = \frac{\Gamma(a)\Gamma(a)}{\Gamma(a+k)\Gamma(a-k)}$$

which implies that $\zeta = 1 \Leftrightarrow k = 0$. That is, none of the moments of X is invariant to the reciprocal transformation.

3. Let $L(\theta, a) = \omega(\theta)(\theta - a)^2$, with $\omega(\theta)$ a non-negative function, be the weighted quadratic loss. Show that $\delta^B(x)$, the estimator that minimizes the Bayesian expected loss $\rho(\pi, \theta)$ has the form

$$\delta^B(x) = \frac{E(\omega(\theta)\theta|x)}{E(\omega(\theta)|x)}.$$

Hint: Show that any other estimator has a larger Bayesian expected loss.

Solution:

Let $\delta(x)$ be any decision rule,

$$\begin{aligned}
\rho(\pi, \delta) &= \int_{\Theta} w(\theta)(\theta - \delta)^2 f(\theta | x) d\theta = \int_{\Theta} w(\theta)(\theta - \delta^B + \delta^B - \delta)^2 f(\theta | x) d\theta \\
&= \int_{\Theta} w(\theta) [(\theta - \delta^B)^2 - 2(\theta - \delta^B)(\delta^B - \delta) + (\delta^B - \delta)^2] f(\theta | x) d\theta \\
&= \int_{\Theta} w(\theta)(\theta - \delta^B)^2 f(\theta | x) d\theta - 2 \int_{\Theta} w(\theta)(\theta - \delta^B)(\delta^B - \delta) f(\theta | x) d\theta \\
&\quad + \int_{\Theta} (\delta^B - \delta)^2 w(\theta) f(\theta | x) d\theta \\
&= \int_{\Theta} w(\theta)(\theta - \delta^B)^2 f(\theta | x) d\theta + \int_{\Theta} (\delta^B - \delta)^2 w(\theta) f(\theta | x) d\theta \\
&\quad - 2(\delta^B - \delta) \left[\int_{\Theta} w(\theta)\theta f(\theta | x) d\theta - \delta^B \int_{\Theta} w(\theta) f(\theta | x) d\theta \right] \\
&= \int_{\Theta} w(\theta)(\theta - \delta^B)^2 f(\theta | x) d\theta + \int_{\Theta} (\delta^B - \delta)^2 w(\theta) f(\theta | x) d\theta \\
&\quad - 2(\delta^B - \delta) [E(\omega(\theta)\theta|x) - \delta^B E(\omega(\theta)|x)] \\
&= \int_{\Theta} w(\theta)(\theta - \delta^B)^2 f(\theta | x) d\theta + \int_{\Theta} (\delta^B - \delta)^2 w(\theta) f(\theta | x) d\theta \\
&\quad - 2(\delta^B - \delta) \left[E(\omega(\theta)\theta|x) - \frac{E(\omega(\theta)\theta|x)}{E(\omega(\theta)|x)} E(\omega(\theta)|x) \right]
\end{aligned}$$

and notice that the first term does not depend on δ . Thus, the rule that minimizes $\rho(\pi, \delta)$ is that which makes the second term equal to zero, namely $\delta = \delta^B$.

4. Consider $x | \theta \sim N(\theta, 1)$, $\theta \sim N(0, 1)$ and the loss

$$L(\theta, a) = e^{3\theta^2/4}(\theta - a)^2.$$

- (a) Show that the estimator that minimizes the Bayesian expected posterior loss in this case is $\delta(x) = 2x$. Hint: use the previous exercise.
(b) Show that $\delta_0(x) = x$ dominates $\delta(x)$.

Solution:

- (a) First notice that if $X | \theta \sim N(\theta, 1)$ and $\theta \sim N(0, 1)$ then $\theta | X \sim N(\frac{x}{2}, \frac{1}{2})$. Now, taking $w(\theta) = e^{3\theta^2/4}$, from the previous exercise is known that $\delta^B(x) = \frac{E(\omega(\theta)\theta|x)}{E(\omega(\theta)|x)}$.

In this case

$$\begin{aligned}
E(w(\theta)\theta \mid x) &= \int_{\mathbf{R}} \theta \exp\left\{\frac{3\theta^2}{4}\right\} \frac{1}{\sqrt{\pi}} \exp\left\{-\left(\theta - \frac{x}{2}\right)^2\right\} d\theta \\
&= \int_{\mathbf{R}} \theta \frac{1}{\sqrt{\pi}} \exp\left\{-\frac{1}{4}\theta^2 + \theta x - \frac{x^2}{4}\right\} d\theta \\
&= \exp\left\{\frac{3x^2}{4}\right\} \int_{\mathbf{R}} \theta \frac{1}{\sqrt{\pi}} \exp\left\{-\frac{1}{4}(\theta^2 - 2\theta x - 4x^2)\right\} d\theta \\
&= 2 \exp\left\{\frac{3x^2}{4}\right\} \int_{\mathbf{R}} \theta \frac{1}{\sqrt{2\pi}2} \exp\left\{-\frac{1}{2 \cdot 2}(\theta - 2x)^2\right\} d\theta \\
&= 2 \exp\left\{\frac{3x^2}{4}\right\} E[\psi] = 4x \exp\left\{\frac{3x^2}{4}\right\}
\end{aligned}$$

where $\psi \sim N(2x, 2)$.

Similarly,

$$\begin{aligned}
E(w(\theta) \mid x) &= 2 \exp\left\{\frac{3x^2}{4}\right\} \int_{\mathbf{R}} \frac{1}{\sqrt{2\pi}2} \exp\left\{-\frac{1}{2 \cdot 2}(\theta - 2x)^2\right\} d\theta \\
&= 2 \exp\left\{\frac{3x^2}{4}\right\}
\end{aligned}$$

Thus, $\delta^B = 2x$.

(b) Notice that $R(\delta, \theta)$ satisfies

$$\begin{aligned}
R(\delta, \theta) &= E(w(\theta)(\theta - 2x)^2) = w(\theta)E(\theta^2 - 4\theta x + 4x^2) \\
&= w(\theta)\{\theta^2 - 4\theta^2 + 4(1 + \theta^2)\} = w(\theta)(\theta^2 + 4).
\end{aligned}$$

while $R(\delta_0, \theta)$ is given by

$$\begin{aligned}
R(\delta_0, \theta) &= E(w(\theta)(\theta - x)^2) = w(\theta)E(\theta^2 - 2\theta x + x^2) \\
&= w(\theta)\{\theta^2 - 2\theta^2 + (1 + \theta^2)\} = w(\theta).
\end{aligned}$$

Thus, $R(\delta_0, \theta) < R(\delta, \theta)$ for all θ , which implies that δ_0 dominates δ .

5. Assume you have to guess a secret number θ . You know that θ is an integer. You can perform an experiment that would yield either the number before it or the number after it, with equal probability. You perform the experiment twice. More formally, let x_1 and x_2 be independent observations from

$$f(x = \theta - 1 \mid \theta) = f(x = \theta + 1 \mid \theta) = 1/2.$$

Consider the 0-1 loss function, i.e.,

$$L(\theta, a) = \begin{cases} 1 & a \neq \theta \\ 0 & \text{otherwise} \end{cases}$$

- (a) Find the risks of the estimators $\delta_0(x_1, x_2) = \frac{x_1 + x_2}{2}$ and $\delta_1(x_1, x_2) = x_1 + 1$.

(b) Find the estimator $\delta^B(x_1, x_2)$ that minimizes the Bayesian expected loss.

Solution:

(a) By the definition,

$$R(\delta_0, \theta) = E \left[L \left(\frac{X_1 + X_2}{2}, \theta \right) \right] = 0P(X_1 \neq X_2) + 1P(X_1 = X_2) = \frac{1}{2}.$$

since X_1 and X_2 are independent. Similarly,

$$R(\delta_1, \theta) = E[L(X_1 + 1, \theta)] = 0P(X_1 = \theta - 1) + 1P(X_1 = \theta + 1) = \frac{1}{2}.$$

(b) To obtain the posterior distribution of θ , we consider the following two cases;

i. if $x_1 \neq x_2$

$$p \left(\theta = \frac{x_1 + x_2}{2} \mid x_1, x_2 \right) = 1 \quad \text{and} \quad p(\theta = \theta_0 \mid x_1, x_2) = 0 \text{ for } \theta_0 \neq \frac{x_1 + x_2}{2}.$$

In words, the posterior probability is a mass point at $\theta = (x_1 + x_2)/2$. Thus, in this case, $\delta^B(X_1, X_2) = \frac{X_1 + X_2}{2}$ since $E(L(\delta^B, \theta) \mid X_1 \neq X_2) = 0$.

ii. if $x_1 = x_2 = x$

$$\begin{aligned} p(\theta = x + 1 \mid x_1, x_2) &= \frac{p(\theta = x + 1 \mid x_1, x_2)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{p(x \mid \theta = x + 1)\pi(\theta = x + 1)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{\frac{1}{2}\pi(\theta = x + 1)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{\frac{1}{2}\pi(\theta = x + 1)}{\frac{1}{2}\pi(\theta = x - 1) + \frac{1}{2}\pi(\theta = x + 1)} \\ &= \frac{\pi(\theta = x + 1)}{\pi(\theta = x - 1) + \pi(\theta = x + 1)} \end{aligned}$$

$$\begin{aligned} p(\theta = x - 1 \mid x_1, x_2) &= \frac{p(\theta = x - 1 \mid x_1, x_2)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{p(x \mid \theta = x - 1)\pi(\theta = x - 1)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{\frac{1}{2}\pi(\theta = x - 1)}{p(\theta = x - 1 \mid x_1, x_2) + p(\theta = x + 1 \mid x_1, x_2)} \\ &= \frac{\frac{1}{2}\pi(\theta = x - 1)}{\frac{1}{2}\pi(\theta = x - 1) + \frac{1}{2}\pi(\theta = x + 1)} \\ &= \frac{\pi(\theta = x - 1)}{\pi(\theta = x - 1) + \pi(\theta = x + 1)} \end{aligned}$$

$$p(\theta = x + j \mid x_1, x_2) = 0 \quad \forall j \in \mathbb{Z} - \{-1, 1\}$$

In words, the posterior supports two values, $\{x + 1, x - 1\}$. Since $\theta - 1$ and $\theta + 1$ are equally likely, the posterior probabilities of θ being $x + 1$ and $x - 1$ are proportional to their prior probabilities.

Consider first $\delta^B(X_1, X_2) = x + 1$. The resulting expected loss is

$$E(L(\delta^B, \theta) \mid x_1 = x_2 = x) = \frac{\pi(\theta = x - 1)}{\pi(\theta = x - 1) + \pi(\theta = x + 1)}.$$

Next, we consider $\delta^B(X_1, X_2) = x - 1$. The resulting expected loss is

$$E(L(\delta^B, \theta) \mid x_1 = x_2 = x) = \frac{\pi(\theta = x + 1)}{\pi(\theta = x - 1) + \pi(\theta = x + 1)}.$$

Therefore, the Bayes estimator is

$$\delta = \begin{cases} \frac{x_1 + x_2}{2} & \text{if } x_1 \neq x_2, \\ x + 1 & \text{if } x_1 = x_2 = x \text{ and } \pi(\theta = x - 1) < \pi(\theta = x + 1) \\ x - 1 & \text{if } x_1 = x_2 = x \text{ and } \pi(\theta = x - 1) \geq \pi(\theta = x + 1). \end{cases}$$

6. Consider a point estimation problem in which you observe x_1, \dots, x_n as i.i.d. random variables of the Poisson distribution with parameter θ . Assume a squared error loss and a prior of the form $\theta \sim \text{Gamma}(\alpha, \beta)$.
 - (a) Show that the Bayes estimator is $\delta^B(x) = a + b\bar{x}$ where $a > 0$, $b \in (0, 1)$ and $\bar{x} = \sum_{i=1}^n x_i / n$. You may use the fact that the distribution of $\sum_i x_i$ is Poisson with parameter θn without proof.
 - (b) Find the MLE for θ (Note: to remind how to find MLEs, read Casella and Berger, Section 7.2.2— see Def 7.2.4).
 - (c) Compute and graph the frequentist risks of $\delta^B(x)$ and that of the MLE.
 - (d) Compute the Bayes risk of $\delta^B(x)$.
 - (e) Suppose that an investigator wants to collect a sample that is large enough that the Bayes risk after the experiment is half of the Bayes risk before the experiment. Find that sample size.

Solution:

- (a) The sufficient statistics, $T = \sum_{i=1}^n X_i \sim \text{Poi}(n\theta)$. So

$$\pi(\theta \mid t) \propto e^{-n\theta} (n\theta)^t \theta^{\alpha-1} e^{-\beta\theta}.$$

That is, $\theta \mid t \sim \text{Gamma}(\alpha + t, \beta + n)$. We know that under squared error loss, the Bayes estimator is the posterior mean, in this case

$$\delta^B(x) = \frac{\alpha + t}{\beta + n} = a + b\bar{x}$$

where $a = \alpha / (\beta + n)$ and $b = n / (\beta + n)$.

(b) $f(\mathbf{x} \mid \theta) \propto e^{-n\theta} \theta^t$ so $\log(f(\mathbf{x} \mid \theta)) \propto -n\theta + t \log(\theta)$.

Then,

$$\frac{\partial \log(f(\mathbf{x} \mid \theta))}{\partial \theta} = -n + \frac{t}{\theta}$$

$$\Rightarrow \hat{\theta} = t/n = \bar{x}.$$

(c) First notice that $E(\bar{x}) = \frac{1}{n}E(t) = \theta$, $\text{Var}(\bar{x}) = \frac{1}{n^2}\text{Var}(t) = \frac{\theta}{n}$ and $E(\bar{x}^2) = \frac{\theta}{n} + \theta^2$. Therefore,

$$R(\hat{\theta}, \theta) = E((\bar{x} - \theta)^2) = E(\bar{x}^2) - 2\theta E(\bar{x}) + \theta^2 = \frac{\theta}{n}$$

while

$$\begin{aligned} R(\delta, \theta) &= E((a + b\bar{x} - \theta)^2) = E(a^2 + 2a(b\bar{x} - \theta) + (b\bar{x} - \theta)^2) \\ &= (b-1)^2\theta^2 + \left[2a(b-1) + \frac{b^2}{n}\right]\theta + a^2 \\ &= \frac{1}{(\beta+n)^2} (\beta^2\theta^2 + (-2\alpha\beta + n)\theta + \alpha^2). \end{aligned}$$

(d) Since $\theta \sim \text{Gamma}(\alpha, \beta)$, we have $E(\theta) = \alpha/\beta$ and $E(\theta^2) = \alpha/\beta^2 + \alpha^2/\beta^2$.

$$\begin{aligned} r(\pi, \delta) &= E \left[\frac{1}{(\beta+n)^2} (\beta^2\theta^2 + (-2\alpha\beta + n)\theta + \alpha^2) \right] \\ &= \frac{\alpha}{\beta(\beta+n)}. \end{aligned}$$

(e) The Bayes risk before the experiment under squared error loss is the prior variance, that is, α/β^2 . We want to find n such that

$$\frac{\alpha}{\beta(\beta+n)} < \frac{\alpha}{2\beta^2}.$$

Now we solve for n and get $n > \beta$.

7. A loss function investigated by Zellner (1986) is the LINEX (LINear-EXponential) loss, a loss function that can handle asymmetries in a smooth way. The LINEX loss is given by

$$L(\theta, a) = e^{c(a-\theta)} - c(a-\theta) - 1,$$

where c is a positive constant. As the constant c varies, the loss function varies from very asymmetric to almost symmetric.

Let X_1, \dots, X_n be iid $N(\theta, \sigma^2)$, where σ^2 is known, and suppose that θ has the noninformative prior, $\pi(\theta) \propto 1$. Show that the Bayes estimator of θ under LINEX loss is given by

$$\delta^B(\bar{X}) = \bar{X} - \frac{c\sigma^2}{2n}.$$

Solution:

In general, for any estimator a

$$\rho(\pi, a) = E[L(\theta, a) \mid X] = E[e^{c(a-\theta)} - c(a-\theta) - 1 \mid X] = e^{ca}E[e^{-c\theta} \mid X] - ca + cE[\theta \mid X] - 1.$$

Then

$$\frac{d\rho(\pi, a)}{da} = ce^{ca}\mathbb{E}[e^{-c\theta} \mid X] - c = 0 \quad \Leftrightarrow \quad a^B = \frac{-\ln(\mathbb{E}[e^{-c\theta} \mid X])}{c}$$

and

$$\frac{d^2\rho(\pi, a)}{da^2} = c^2e^{ca}\mathbb{E}[e^{-c\theta} \mid X] > 0$$

so, a^B is a minimizer.

Now, in particular, $\theta \mid \mathbf{X} \sim \mathcal{N}\left(\bar{x}, \frac{\sigma^2}{n}\right)$ which implies that $e^{-c\theta} \mid \mathbf{X} \sim \text{LogNormal}\left(-c\bar{x}, \frac{c^2\sigma^2}{n}\right)$.

Therefore, $\mathbb{E}[e^{-c\theta} \mid X] = e^{-c\bar{x} + \frac{c^2\sigma^2}{2n}}$ and $a^B = \bar{x} - \frac{c\sigma^2}{2n}$ as desired.